

Machine Learning Basics

Gradient descent minimizes the loss function by iteratively adjusting weights using the gradient. Learning rate controls step size. Backpropagation applies the chain rule to compute gradients across layers. Overfitting occurs when a model memorizes training data; regularization (L1/L2) and dropout help prevent it. Bias-variance tradeoff: high bias = underfitting, high variance = overfitting.